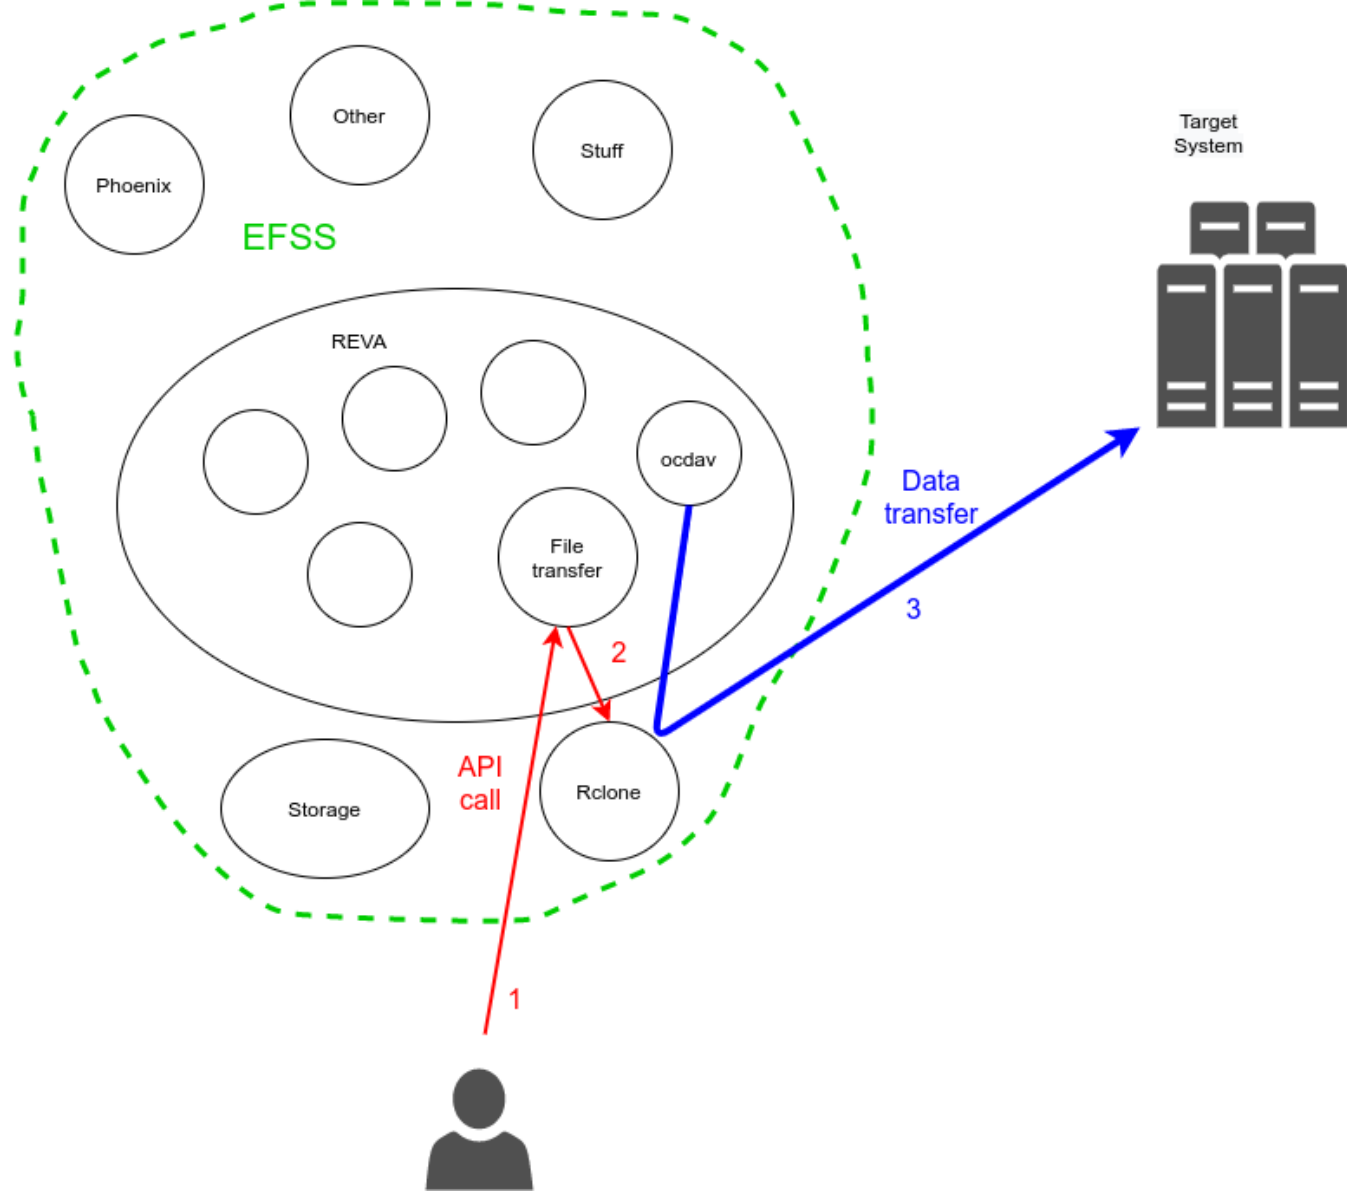


Software Design Rclone Service

Remarks

- Still work in progress
- This is just to get us going
- Lay the ground work for later development
 - Start with rclone, later on FileSender, FTS, Rucio



Rclone

- Client that can copy data to and from just about any cloud storage known to man amongst many other things
- Can run as a daemon “rclone rcd” accepting http POST requests in JSON
- Needs a config file to define “remotes”
- Can’t do third-party copies (yet), i.e. data traverse through the rclone process

Context

- Data locality
- Service can be part of REVA, but also be standalone → multivendor
- Start with rclone.
 - Nothing on the critical path
 - Pathfinder for the FileSender/FTS/Rucio work
- REVA is about data/application sharing.
 - This application should run at every EFSS system so it won't be shared

Goals

- HTTP service that can perform “rclone copy” commands
 - Design REST API
 - Make sure it can be re-used for FileSender/FTS/Rucio
 - Front-end to accept http requests
 - Backend to do the data transfer heavy lifting

Non-Goals

- Implementing any of the other rclone commands
- Service accessible over WAN
- Integration with web I/F of EFSS solutions
- High Availability (retry mechanisms?)
- Data transfer management
- Scalability (need to look at this later on, though)

Existing Solutions

- Do your own scripting to copy 200TB of data to a EFSS (yuck)

Proposed Solution

- EFSS systems in general not very well suited for handling the larger data sets
- HTTP-based service to do this dirty work

Alternative Solutions

- Nope

Testing, Monitoring Alerting

- TODO

Cross-Team Impact

- None
- This is very much a stand-alone thingie which has very little impact on the rest of the CS3MESH project

Open Questions

- Security
- Scalability
 - For filesender/FTS/Rucio this may be less of an issue but since data passes through rclone this should be taken into account for this case
- High-Availability
- Multi-user issues

Nitty Gritty Details

Basic workflow

- 1) Request comes in
- 2) Process is forked to handle new incoming requests
- 3) Parent process validates the input
 - 1) Security
 - 2) Validity
- 4) Rclone daemon is fired up “rclone rcd”
- 5) Issue “POST config” command to it to create config file in memory to create the necessary “remotes”
- 6) Issue “POST copy” command to do the transfer
- 7) After the transfers have ended, the stuff that rclone returns (exit code, errors, whatever) is returned to the client issuing the request in step 1
- 8) Cleanup: Delete rclone config, rclone daemon is killed, process is killed

Timeline

- Coming soon in theatres near you
- Start the development work
- Determine what other partners can/want to contribute
- Look ahead to FileSender/FTS/Rucio